

SAGE: New Math for Next Generation AI Infrastructure



Most AI portfolios are powered by vastly inferior math.

SAGE MPW provides superior math — Training at $O(n)$. Inference at $O(n \log n)$.

01 AI breakthrough — it's just math

Quadratic cost drops to linear for training and near-linear for inference.

TRAINING

$$\cancel{O(n^2)} \rightarrow O(n)$$

INFERENCE

$$\cancel{O(n^2)} \rightarrow O(n \log n)$$

02 Market implications

a ~~GW~~ → **kW**

Kilowatts instead of gigawatts

Data-center power and capital requirements fall by orders of magnitude — greener compute, with planetary-scale benefits.

“The best data center is the one you do not need to build.”

— Paul Charlton

b **Train + infer** on battery

Battery-powered devices do both in near real time

Personal and sovereign AI: own your inference instead of renting it from untrustworthy partners.

Smartphones

Smart infrastructure

03 Track record — who you are betting on

Paul Charlton Founder & tech lead
Inevitable Systems

Three decades building the hidden system layers that let computing cross its hardest thresholds: ubiquitous digital video, Java across devices, routers across cores, filesystems across operating systems, institutions across cloud, global-scale IoT.

Apple QuickTime · Sun Java2D · Cisco IOS multicore · Intuit QuickBooks · SGI CXFS · HP storage · D.E. Shaw Cloud

EVIDENCE

Digital Video — three decades of ubiquity

youtube.com/watch?v=41DLFG2ohzk&t=2774s

Team Leader — Paul Charlton

linkedin.com/in/techguru

Related portfolio

research.buiq.org


— THE OLD MATH

The problem: everything is multiplied by everything.

40 85 79 26 57 87
70 43 80 39 34 70
39 91 29 76 59 11


For two centuries, one move has run our math. [Unsustainable...](#)

Matrix multiplication became the default engine of AI.




[Multiplying everything by everything to find the little that matters is the most expensive way to compute a statistic.](#)

Every bottleneck is a downstream cost of one upstream choice.



A few hundred places on Earth.

Where frontier intelligence can exist.
[Everyone else rents it by the token.](#)



— THE NEW MATH

SAGE.

Statistical Algebra with Grounded Evidence.

Integer-only

Nonparametric

INTEGERS

412

7

1,286

53

9

EVIDENCE

Run the statistics in integers, directly over the evidence, and keep the evidence.

No matrices to multiply. No backpropagation. No GPU.

Precise math can be small.
Small can run anywhere.
Anywhere means owned, not rented.
Evidence means provable.



- ✓ Every answer traces to its evidence.
- ✓ Any record can be erased on demand.
- ✓ Trains and runs on a CPU, from a Raspberry Pi to a rack.

Company claims, pending Reveal demonstration.

STATUS TODAY Internal prototype. Not yet validated at scale on comparison hardware.

By end of year: the Reveal comparison demonstration, on a Raspberry Pi and in the data center, with an independent reviewer.


Intelligence. Unlocked.

More intelligence per watt. Lighter on the Earth.
Knowledge belongs to the people who make it.



— THE VALUE

Vertical path to market/valuation.



The VC staircase. Many small steps, each gated by a new round.

Our vertical ladder. One straight climb: \$0 to >\$100B in months.

Company thesis

\$27T¹

AI total addressable market

\$7T²

Data-center capex by 2030

1,000,000×³

Up to 1,000,000× more compute per token per watt.
Same quality. 100× lower total cost of ownership (TCO).
Speedup scales as $\sqrt{\text{parameters}}$ — $10^6\times$ at 1T. No matmul, no GPU.
Even 10% of company-projected scale metrics would eclipse existing matmul without SAGE.

SCENARIO 1 · BIDDING WAR

>\$100B⁴

Value to prospective bidders

Company estimate — anticipated bank-led road show with Reveal. Confidential license bidding, post-Reveal. No sale of the IP. Nvidia · OpenAI · Google · Anthropic · SpaceX · xAI
\$60B Cursor → SpaceX
\$32B Wiz → Google
\$25B CyberArk → Palo Alto
\$20B Groq → Nvidia

SCENARIO 2 · LOW-COST INFERENCE

\$196B⁵

Value of the token flow

167T tokens / month captured
× \$5.00 per M tokens — under half of market
× 12 months
= **\$10B free cash flow / year**
÷ 5.1% cost of funds, indexed to 10-yr Treasury

Undercut all market leaders; capture outsized margin.

SCENARIO 3 · LICENSE STREAM

\$262B⁶

Value of the license stream

Greater of 10% gross token revenue or 1.5× our floor price. Floor = 1/10,000 market price.

Google \$2.85B / yr
OpenAI \$4.0B / yr
Anthropic \$6.5B / yr
= **\$13.4B license fees / year**
÷ 5.1% cost of funds
Token market grows 21× by 2030.

3.2Q

Tokens / month — Google alone
≈ 57% of 5.6Q market

\$10B free cash flow needs 3% of today's total monthly token market — 0.14% of 2030's.

120Q

Tokens / month by 2030, agentic

\$11.25

Per million tokens, blended frontier

The Reveal is proof for the road ahead: bidding, selling inference, and top sector first partners.

We have early, active edge customer licensing opportunities in IoT/telecom, data centers, sovereign deployment, and robotics.

Three scenarios, not forecasts. Scenario values are company estimates.

1 SpaceX S-1, May 2026 (\$26.5T AI). 2 McKinsey, "The cost of compute" (\$6.7T). 3 Company claim.


4 Company estimate; bank-led road show. Comps: SpaceX S-1; 8-K, Aug 14, 2026; company comp set.

5 Company calc. Google I/O, May 2026; Goldman Sachs, May 2026; Axis LLMflation Index, July 2026 (\$11.25); 10-yr Treasury, Sept 24, 2026.

6 Company calc; revenue as token-revenue proxy. Google: tokens × \$2.50/M avg (Ramp, via Artefact). OpenAI and Anthropic run rates, Jul 2026 (Bloomberg). Floor: \$11.25/M ÷ 10,000.

— THE REVEAL

One event. Three paths.



The Reveal is where investor value is created.

Why public? Every bidder's clock starts at once — no slow-walking.

With vs. without SAGE, to Raspberry Pi, to live data center. Metrics documented. Public results. Third party reviewed.

Side-by-side power meters, live

Same corpus, same hardware

Audience types its own prompt

[Independent reviewer](#)

— BIGGEST RISKS

Appropriation

Trade secrets, key patents, exclusive licenses; fast Reveal cuts time-to-market.

Government/Regulatory changes

Non-US IP trust holding entity (for IP protection); aligned with bidders; affiliated lobbying firms.

Founder well-being

No bus crashes: key-person insurance refunds SAFE; in-house succession.

Legit competitors

Each rival splits TAM linearly — room for the top 2–3.

Risks priced in. So why ask for so little?

“You’re only raising \$5M when everyone else asks hundreds of millions — you can’t be real.”

Answer: SAGE doesn’t require the massive training hurdle or months to get there. We’re lean.

— THE RAISE

\$5M at \$50M pre. Same dilution, 1/20 the check.

Others ~\$100M seeds at ~\$1B pre — mostly training runs

Us  \$5M at \$50M pre

Others convert \$100Ms of investor cash into pure entropy — data-center heat.

We train on CPUs. 100× lower total cost of ownership.

Company claim

- 01 The small raise is consistent with the cost claim.
- 02 \$5M buys: Reveal in 2–3 months · two entities + license legal · 6–9 months runway
- 03 Small elite team. Proof in months, not years.
- 04 Witnessed by: bidders at the Reveal; lead-bank diligence before road show.
- 05 Founder track record: Paul Charlton — see our one-pager.

Raise covers cost. Valuation reflects the market.

Why \$50M pre? Priced before proof — your cash finishes it.

\$1.51M funds the Reveal demonstration and key IP. The balance is post-Reveal spend: bidding and inference-scale readiness.

— THE TEAM

Small elite team. Senior, shipped, exited.

Ready to onboard with funds availability.

4

exits

3

ex-Apple

1

Nvidia director

FOUNDER Paul Charlton

Background in our one-pager

Apple: QuickTime on Windows, 1994 · Sun: Java2D glyph rendering, 40× faster. Code still in Android.

D.E. Shaw: gradient descent and SVD over million-dimension spaces, 1990s · IoT at scale: 100M devices, 1T events per day

Principal Engineer

ex-Apple · ex-GoPro

CTO @ 2× exited startups

Current Director of Engineering, Nvidia

Field Engineer

ex-Apple

Chief Architect

Current Fractional CTO · ex-Apple · ex-IdeaLab

CTO, exited startup

Current Principal Engineer, Google

Principal Engineer, exited startup

IoT at scale · ML operations

You are seeding the proof of the new substrate
that unlocks intelligence at the source. For everyone.

We ask for little — then prove it in public.

— NEXT STEP

Standard SAFE angel round, capped, no discount. \$50M pre · \$55M post
Advisory role (active) would be welcome.


USE OF FUNDS

Use of Funds


Oct 2026 – Apr 2027

\$1.51M of SAFE funds the Reveal demonstration and key IP.
The balance is post-Reveal spend: bidding and inference-scale readiness.
Banker / auction / transaction fees outside the \$5M SAFE

- STRUCTURE
- SAGE IP held in an offshore trust for protection: no nation, company or person can shut it down
 - Inevitable Systems, Inc. holds the exclusive master right to license, pre-bid
 - SAFE holders convert at the next round, at IP licensing events
 - No selling the IP outright



BURN BY PHASE



WHERE THE \$5M GOES

Total CapEx	incl. 4 SAGE Racks	\$583K
Total People	Engineering, Exec	\$1.10M
Total OpEx	PR/Production, Recurring, Ins, Travel, Contract Controller	\$667K
IP Protection and Legal	Patent work + initial licensing readiness	\$975K
Reserves / Contingency	Contingency + unallocated reserve	\$1.67M
TOTAL		\$5.0M

[JOIN US](#)



SAGE: unleashing AI for everyone.

One-pager and video

Albreakthrough.ai

Join us

info@inevitable-systems.com

Paul Charlton — portfolio

research.byyiq.org

Paul Charlton — technical CV

techguru.byyiq.org
